

# GeneRic Autonomic Signaling Protocol (GRASP) A Tutorial

Brian Carpenter

October 2021

# Note

- This tutorial is intended for self-study, not for classroom presentation. Please read it at your own speed.
- There are about 60 slides.
- For general background, before continuing, you are strongly advised to read this article first:

[Autonomic Networking Gets Serious](#), Internet Protocol Journal 24(3), pp2-18, October 2021.

# Topics

- Background
- Requirements for an autonomic signaling protocol
- GRASP basics
- GRASP “objectives”
- GRASP operations and messages
- GRASP API
- Logic flows
- Security
- Prototype implementation

# Background and General Terminology

- Autonomic Network: Self-managing (self-configuring, self-protecting, self-healing, self-optimizing) with high-level guidance by a central entity (e.g. the Network Operations Center, NOC)
- Autonomic Function: A specific self-managing feature or function.
- Autonomic Service Agent (ASA): An agent that implements an autonomic function, centralized or distributed.
- Autonomic Node: A node that contains at least one ASA and employs autonomic functions
- Autonomic Control Plane (ACP): Self-configuring fully secure virtual network used for all autonomic messaging.

# Technical Reading List

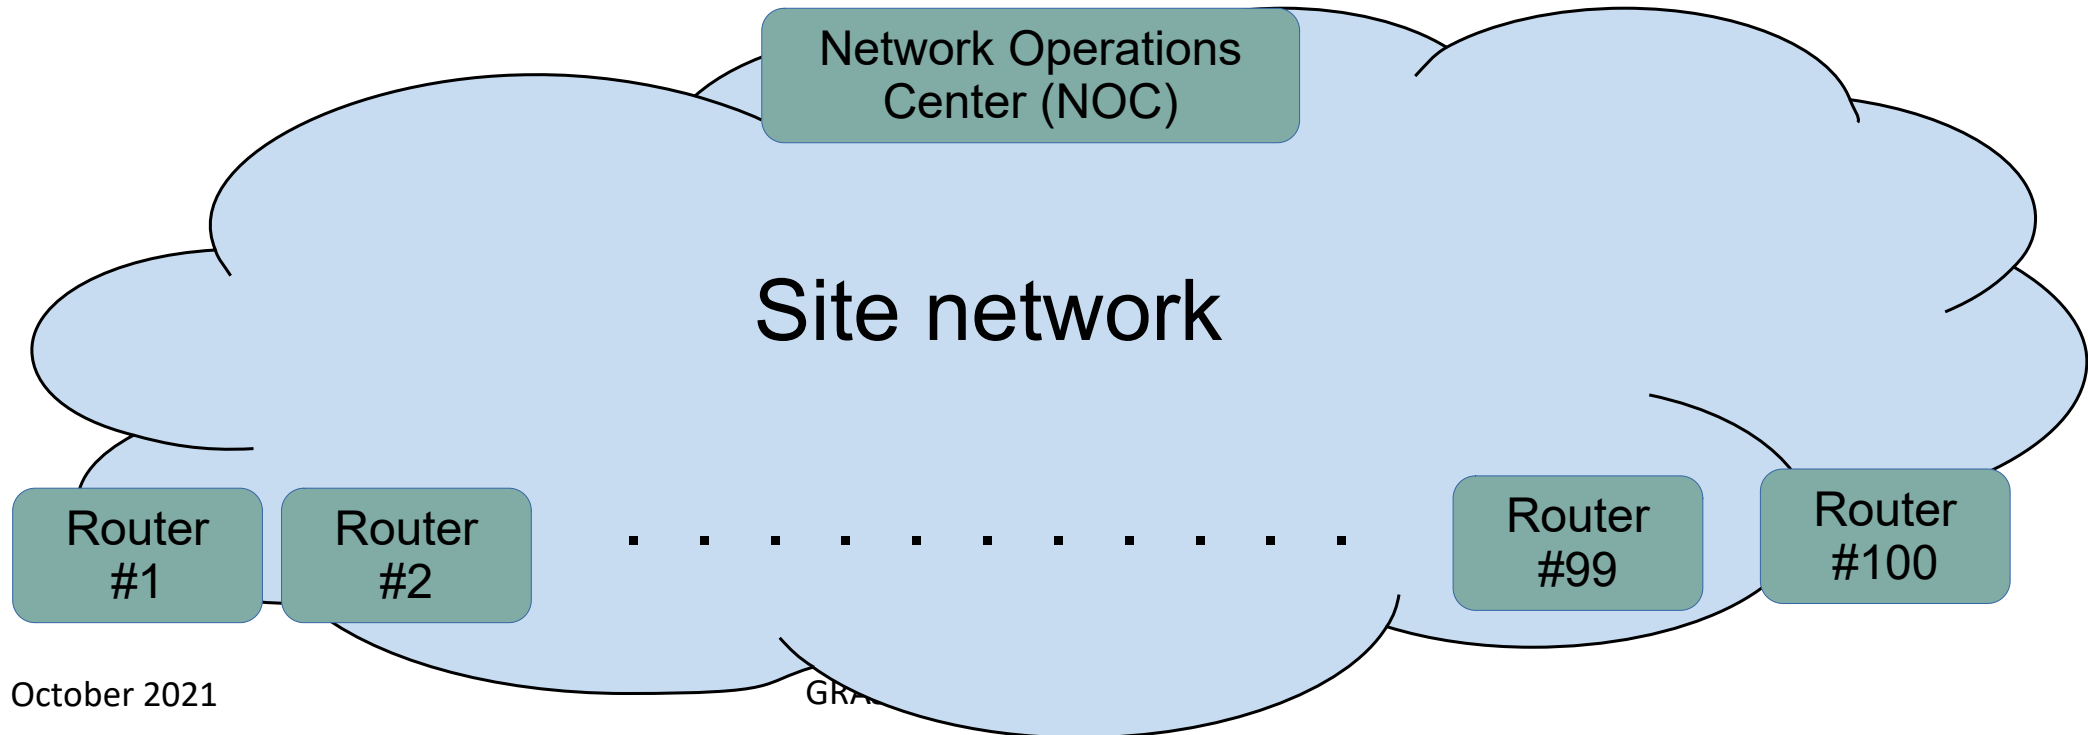
*For all technical details, consult the RFCs:*

- Background: [RFC7575](#) and [RFC8993](#)
- Specifications
  - Autonomic security: [RFC8994](#) and [RFC8995](#)
  - GRASP itself: [RFC8990](#) (protocol) and [RFC8991](#) (API)
- Use cases:
  - Integration with NOC: [RFC8368](#)
  - IP Prefix management: [RFC8992](#)

# Interlude (1)

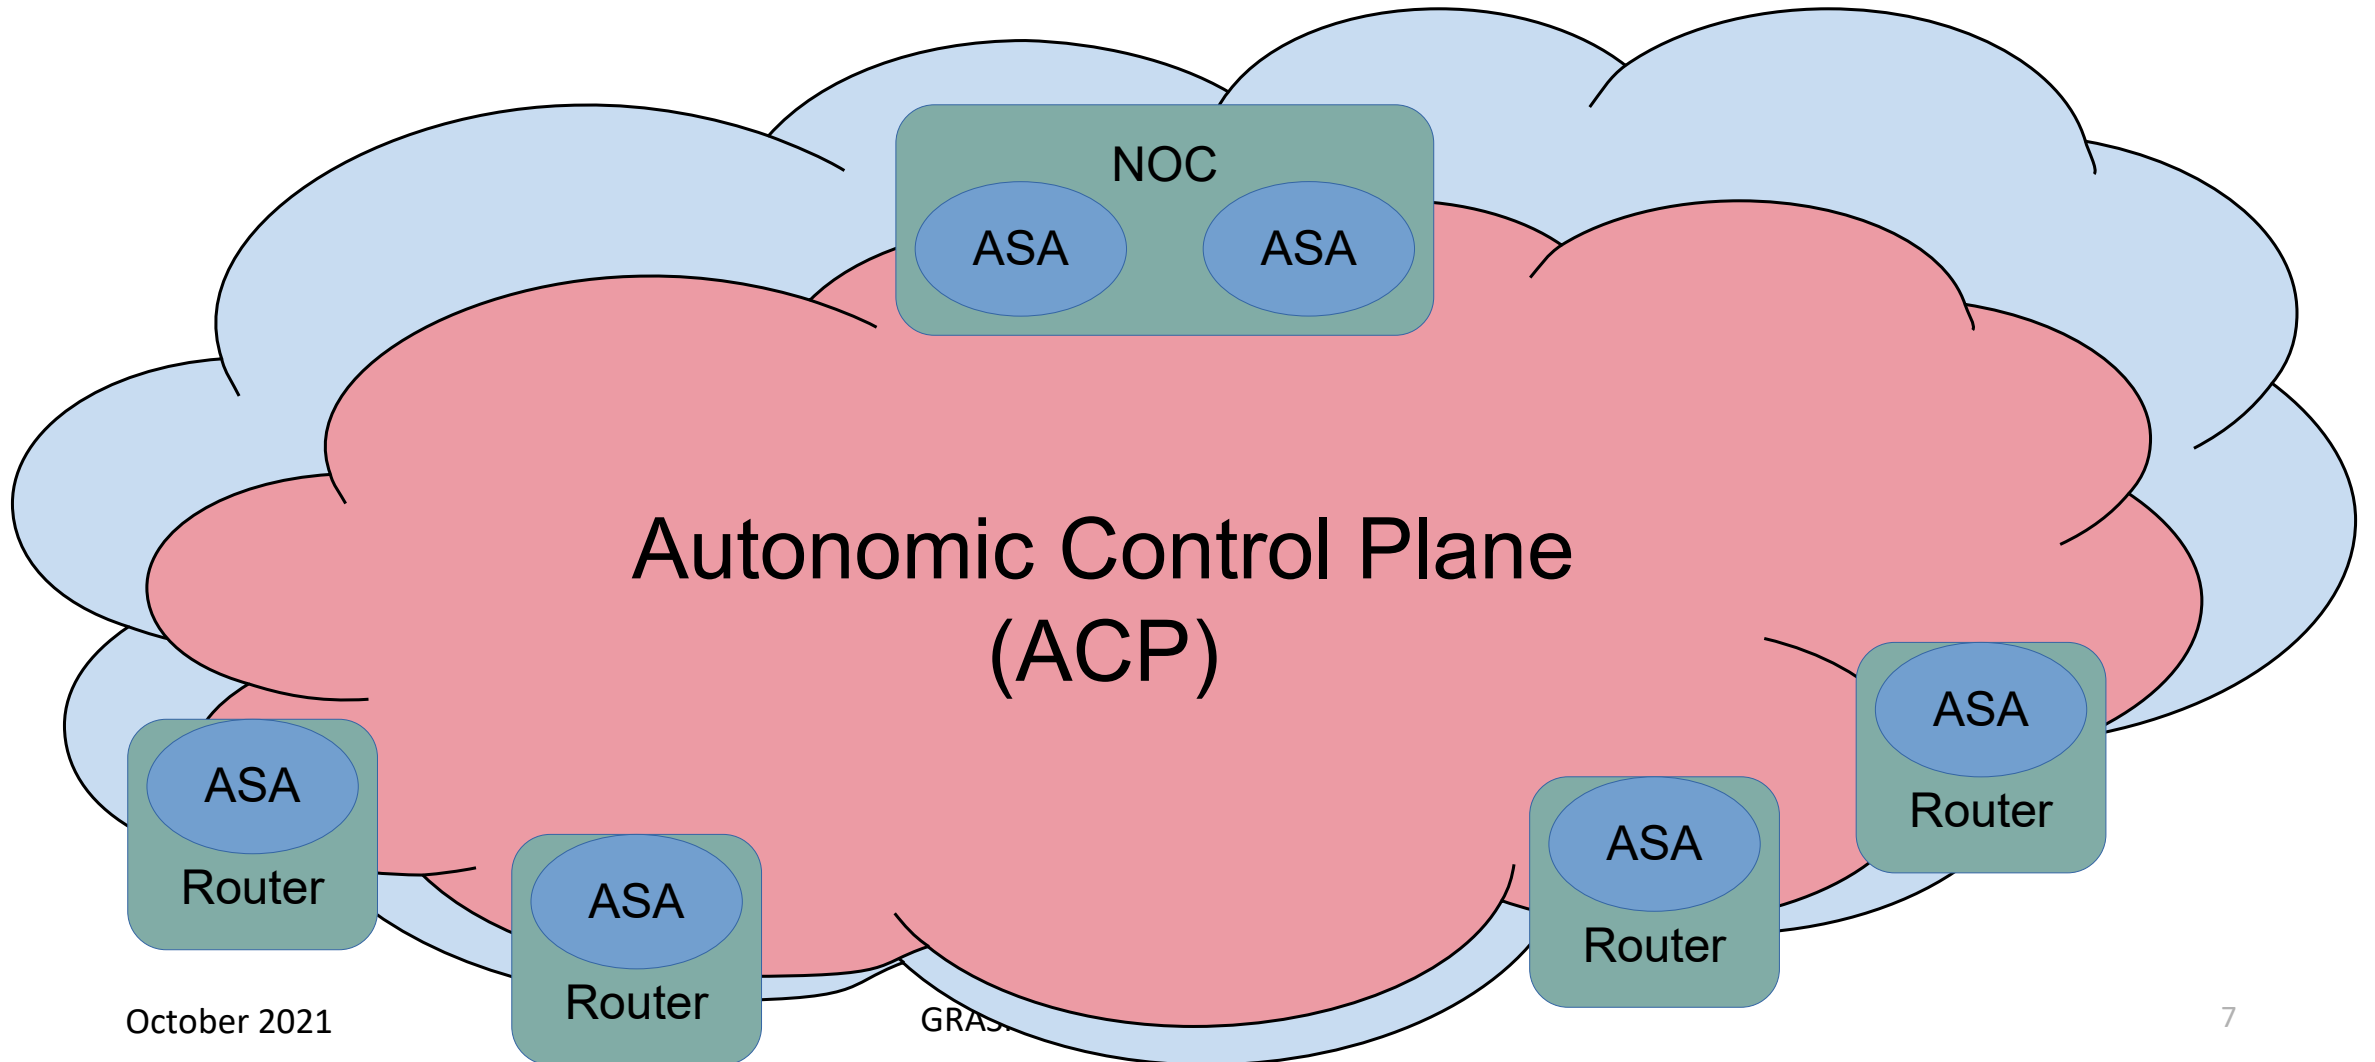


- Imagine a network with a few thousand hosts and (say) a hundred locations. Some resource (e.g., IP address space) must be shared among the locations.
- Can we manage that resource without manual actions?



## Interlude (2)

- Add the ACP and some Autonomic Service Agents



# Requirements for Autonomic Service Agents (ASAs)

- To act autonomically, ASAs need to communicate with each other.
- Even if policy and resources come from a central origin (the NOC), ASAs need to communicate *peer-to-peer*, especially if the network partitions due to an outage.
- There are two primary forms of communication:
  - Configuration or resource data *sent in one direction* only (from one ASA to others);
  - Configuration or resource data *negotiated* between ASAs.



# Why invent a new protocol?

- [RFC8990](#) gives more detail on the ASA communication requirements.
- It also analyzes various existing protocols against these requirements.
- None of them matched up. The IETF ANIMA WG decided that a purpose-designed protocol was the way to go.

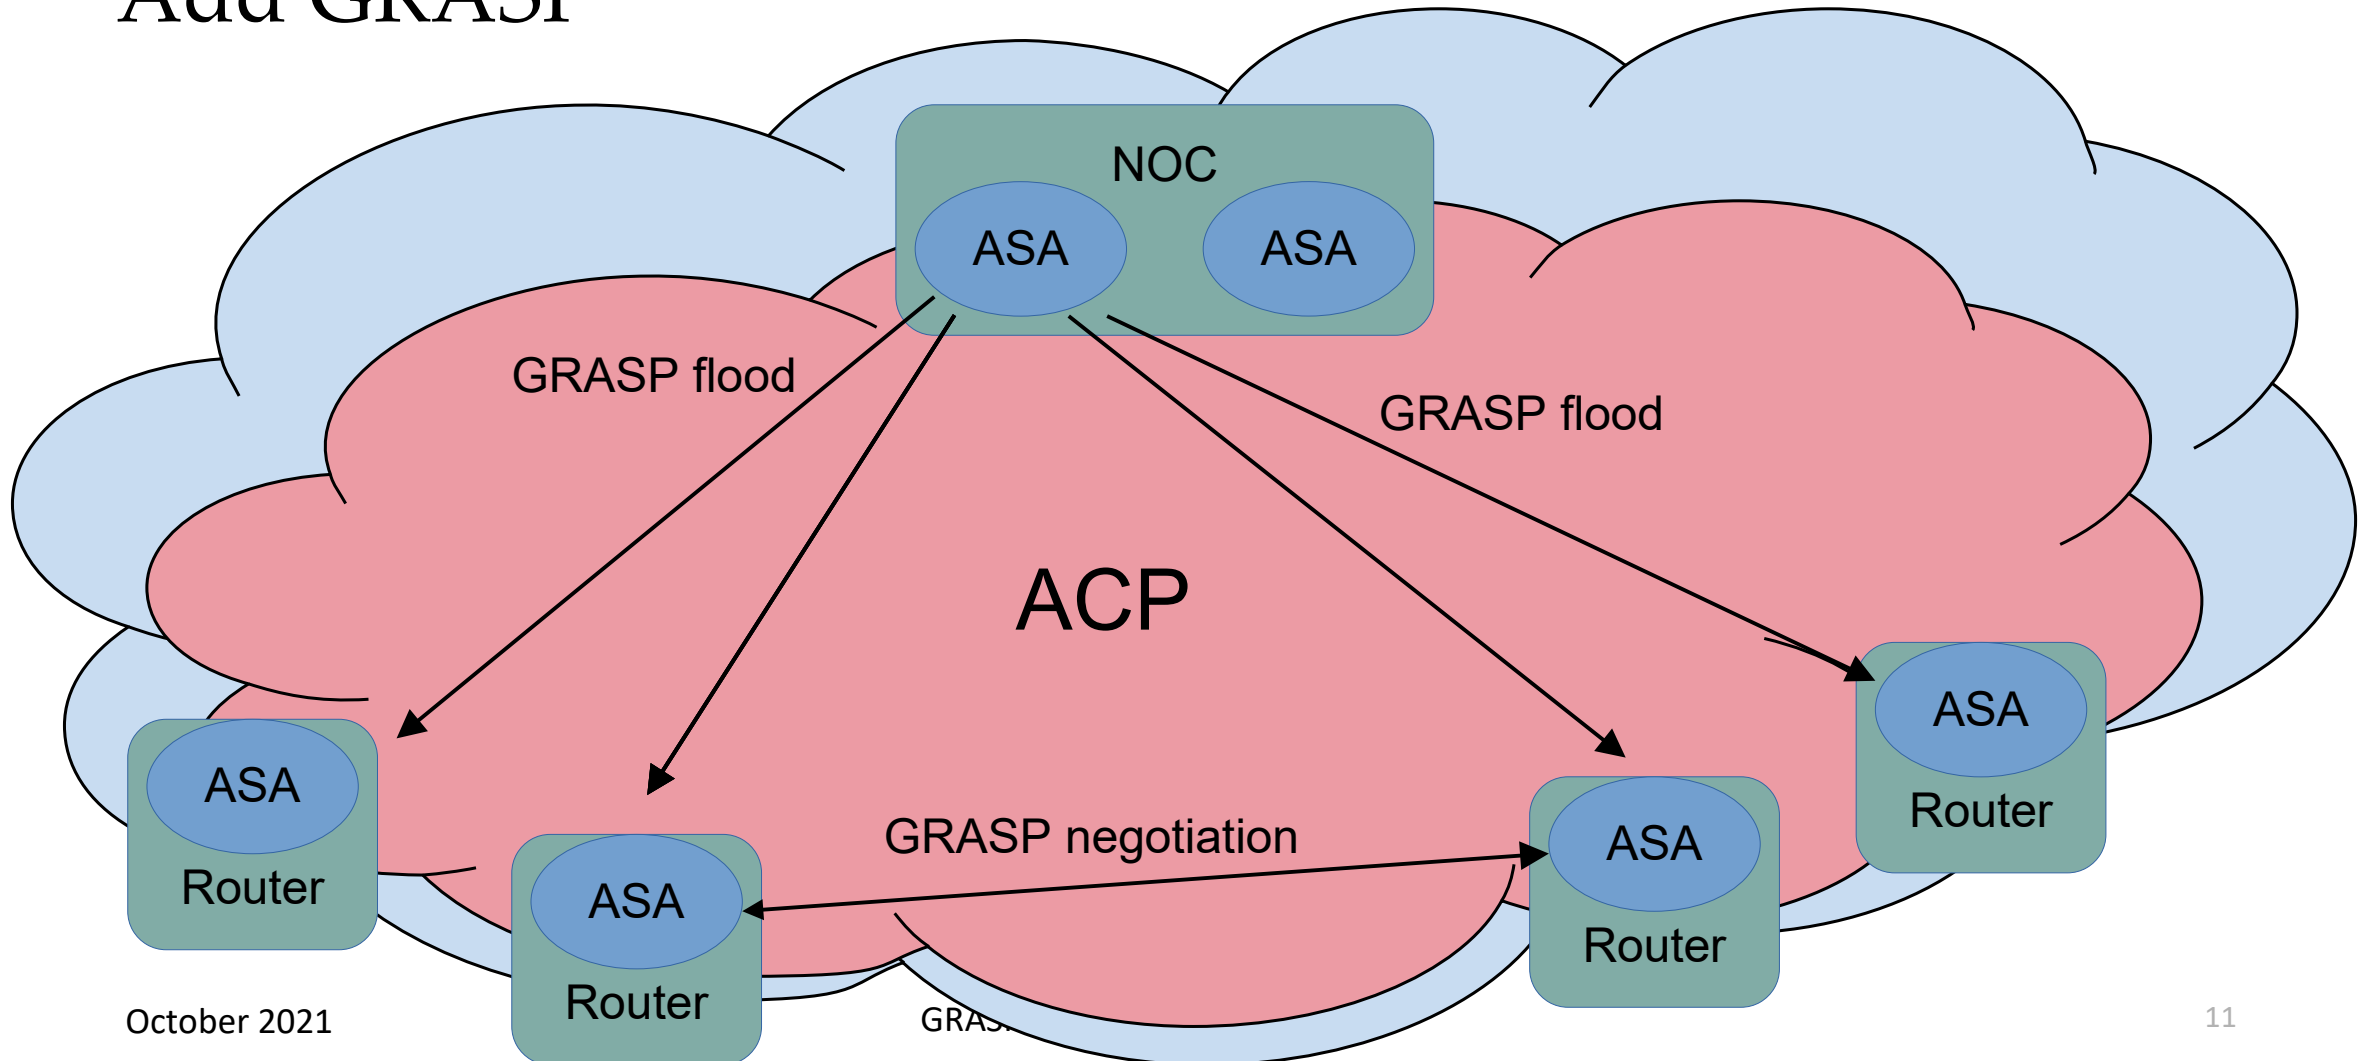
# Basics of GeneRic Autonomic Signaling Protocol (GRASP)

- GRASP will be used for signaling between ASAs
  - That includes the special-purpose ASAs that support secure bootstrap & Autonomic Control Plane (ACP) creation
  - After that, GRASP runs over the ACP to guarantee security
- GRASP provides discovery, flooding, synchronization and negotiation for the technical “objectives” supported by ASAs
  - Based on CBOR (Concise Binary Object Representation)
  - Objectives can be expressed in JSON or Python-like syntax & semantics

# Interlude (3)



- Add GRASP



# Why CBOR? Why not TLV?

- The earliest design used a traditional type-length-value format.
- However, switching to CBOR<sup>1</sup> ([RFC8949](#)) gave
  - much greater flexibility and extensibility,
  - no performance loss,
  - clear protocol definitions in CDDL<sup>2</sup> ([RFC8610](#)),
  - easy mapping to JSON when useful,
  - easy mapping to modern languages like rust and Python.

<sup>1</sup>Concise Binary Object Representation

<sup>2</sup>Concise Data Definition Language

# GRASP Technical Objectives

- In GRASP, an *objective* is a configurable parameter:
  - a logical, numerical or string value, or a more complex data structure.
  - used in Discovery, Negotiation, Flooding and Synchronization.
  - semantics depend on the autonomic function concerned, and are built into the code of each ASA.
- Example for IP prefix management:

```
["PrefixManager", flags, loop_count,  
                             [IP_version, prefix_length, prefix]]
```

# Formal syntax of a GRASP Objective (1)

**objective = [objective-name, objective-flags,  
loop-count, ?objective-value]**

**objective-name = text**

**objective-value = any**

**loop-count = 0..255**

# Formal syntax of a GRASP Objective (2)

**objective-name = text**

This is any human-readable UTF-8 string.

- Generic objectives **MUST NOT** include a colon (":") and **MUST** be registered with IANA.
- Privately defined objectives **MUST** include at least one colon (":"). The string preceding the last colon in the name **MUST** be globally unique, such as a fully-qualified DNS name.

# Formal syntax of a GRASP Objective (3)

## **objective-flags**

A byte containing up to 8 flag bits. Bit numbers defined so far:

**F\_DISC: 0 ; valid for discovery**

**F\_NEG: 1 ; valid for negotiation**

**F\_SYNCH: 2 ; valid for synchronization**

**F\_NEG\_DRY: 3 ; negotiation is a dry run**

F\_NEG and F\_SYNCH cannot both be set to 1.



# Formal syntax of a GRASP Objective (4)

**objective-value = any**

GRASP does not restrict the value field of an objective. Anything that can be expressed in CBOR can be used: for example, a single integer, a UTF-8 string, an array of floating point numbers, or any kind of JSON-like object.

In other words, whatever suits the configuration or optimization task of an ASA is OK.

The specification of a given GRASP objective must define the format of the value field, preferably using CDDL for clarity.

# Formal syntax of a GRASP Objective (5)

**loop-count = 0..255**

In a discovery operation, this variable is used to limit the scope of discovery (see later).

In a negotiable objective ( $F\_NEG = 1$ ), this variable counts down at each step of a negotiation, and the negotiation fails if it reaches zero.

In a synchronizable objective ( $F\_SYNCH = 1$ ), this variable is used to limit the scope of a flooding operation (see later).

## A brown silhouette of a coffee cup with three wavy lines above it representing steam. The cup sits on a matching saucer.

- The flag byte:  
F\_DISC and  
F\_NEG set

# The loop count

# Prefix length

## The actual binary prefix

# GRASP Operations and Messages

- **Discovery** is used by any ASA that needs to discover another (peer) ASA that supports a given objective. It returns zero, one or more responses.
- **Negotiation** is used between two ASAs that support a given objective. They swap values of the objective until negotiation succeeds or fails.
  - GRASP does not support multiparty negotiation; negotiation is 1-to-1.
- **Synchronization** is used between any pair of ASAs that support a given objective. One of them obtains a value of the objective from the other.
- **Flood Synchronization** is used when one ASA needs to distribute the value of a given objective to all others.
  - GRASP does not currently support selective distribution.

The next slide lists the message types that support these operations.

# GRASP Message Types

Discovery (multicast)	M_DISCOVERY
Discovery Response	M_RESPONSE
Request Negotiation	M_REQ_NEG
Negotiation	M_NEGOTIATE
Confirm Waiting	M_WAIT
Negotiation End	M_END
Request Synchronization	M_REQ_SYN
Synchronization	M_SYNCH
Flood Synchronization (multicast)	M_FLOOD
Invalid	M_INVALID
No operation	M_NOOP

# Transport and IP Layer Usage

- Multicasts are IPv6 link-local to port 7017.
  - All modern operating systems support link-local IPv6 by default; nothing to configure.
  - When necessary, GRASP nodes relay these multicasts on other links.
- Unicast messages use a reliable transport protocol over IPv6.
  - Depending on the security environment provided by the ACP, this may be TLS.
  - Otherwise, TCP.
- In a deployment with the standard ACP, the IPv6 environment is self-creating, using Unique Local Addresses, with no operator intervention.

# Formal syntax of a GRASP message (1)

**message-structure = [MESSAGE\_TYPE, session-id,  
                                  ?initiator,  
                                  \*grasp-option]**

**MESSAGE\_TYPE = 0..255**

**session-id = 0..4294967295 ; up to 32 bits**

**grasp-option = any**

## Formal syntax of a GRASP message (2)

**MESSAGE\_TYPE = 0..255**

Just an integer defining the message type.

**session-id = 0..4294967295 ; up to 32 bits**

A unique pseudo-random number identifying a session (discovery, negotiation, etc.). Together with the IP address of the **initiator** of a session, this forms a unique handle, to distinguish simultaneous sessions.



# Formal syntax of a GRASP message (3)

**grasp-option = any**

Defined GRASP options *include*:

**objective** (as above)

**ipv6-locator-option = [O\_IPv6\_LOCATOR, ipv6-address,  
transport-proto, port-number]**

(used in M\_RESPONSE and elsewhere)

**accept-option = [O\_ACCEPT]**

**decline-option = [O\_DECLINE, ?reason]**

(used in M\_END to finish a negotiation)

# Interlude (5)



- M\_NEGOTIATE message in detail

Message type

Session  
identifier

```
[5, 1763418432,  
 ['PrefixManager', 3, 6,  
  [6, 57, b'20010db8ffffffff8000000000000000000000']]]
```

GRASP objective  
as before

## Interlude (6)

- CBOR as it appears on the wire

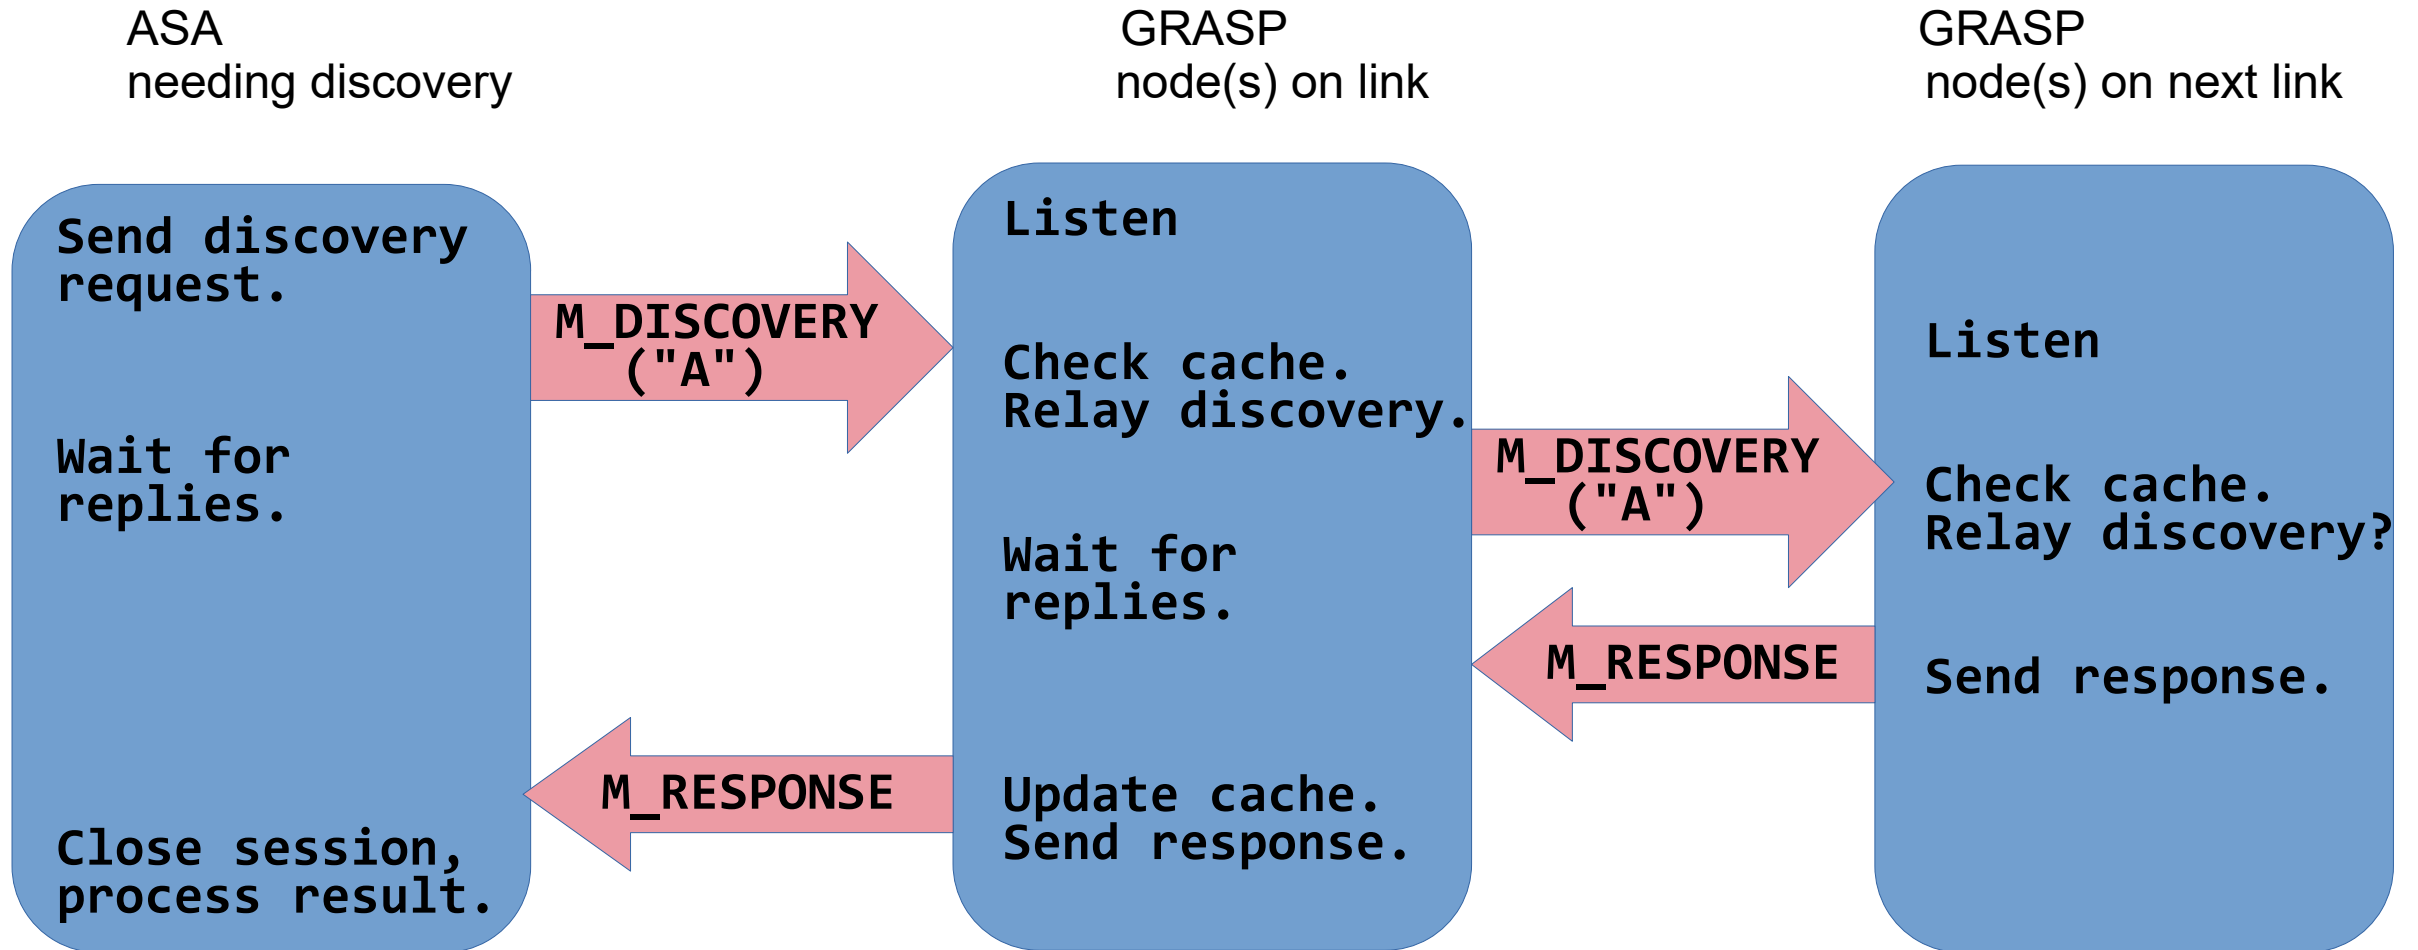
```
b'4df6aa69862ca8d0fc0b2dcd88c726888bf905b7a2fb3a  
8ba501606e35932dadb7474a3ef5c43c3a9a5e58fbcbbc95  
1d'
```

- 48 bytes

# What happens during Discovery

- An ASA that needs to find a peer handling objective “A” originates an M\_DISCOVER message for objective “A”.
  - This goes out as a link-local UDP multicast on each of the node’s interfaces to the ACP.
- Every GRASP node that receives the multicast and supports “A” or has cached the address of a node that supports “A” sends a unicast M\_RESPONSE back, including an **ipv6-locator-option**.
- Any node which is also a router to other links will relay the M\_DISCOVER to its other interfaces, and then relay back any M\_RESPONSEs that it receives (and cache them).
- The original node will return all results received before a specified timeout to the requesting ASA (and cache them).

# A discovery process

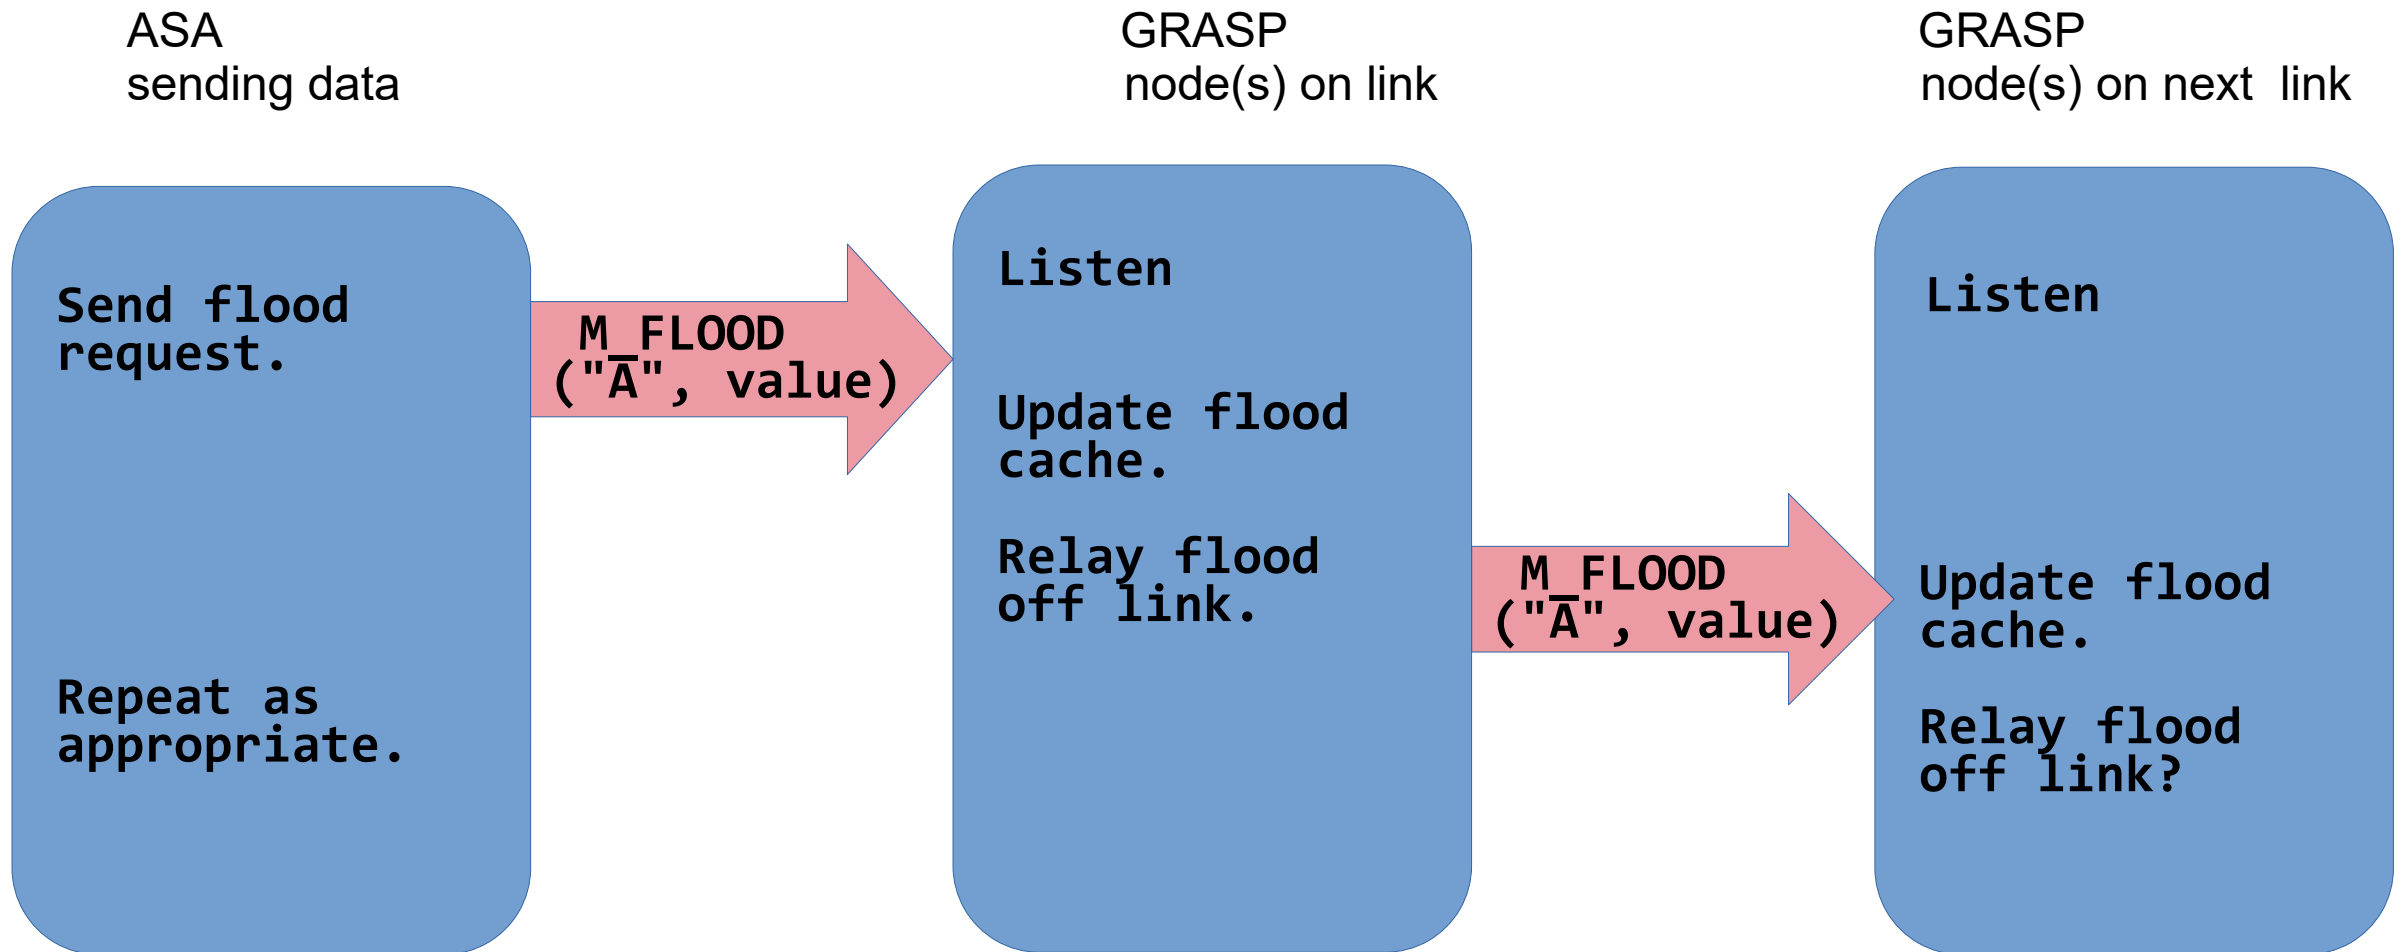


(Relaying is limited by loop count.)

# What happens during Flooding

- An ASA handling objective “A” originates an M\_FLOOD message for objective “A”, including whatever value of the objective it wants to send to all nodes. This goes out as a link-local UDP multicast on each of the node’s interfaces to the ACP.
- Every GRASP node that receives the multicast caches a copy of the objective and its value for local use.
- Any node which is also a router to a different physical link will also relay the M\_FLOOD to its other interfaces.
- Floods may specify an expiry timeout, after which other nodes will mark the cached value as expired.

# A flooding process



(Relaying is limited by loop count.)

# What happens during Synchronization

- An ASA able to provide a value for objective “A” waits for any incoming M\_REQ\_SYN message. This also makes the ASA discoverable for “A”.
  - This is server-like behavior
- An ASA needing to obtain a value for “A” first uses discovery to find all peers that support “A” and chooses one of them. It then originates a unicast M\_REQ\_SYN to the chosen peer.
- The peer responds with a unicast M\_SYNCHRONIZE message containing the current value of “A”.

(A diagram for synchronization comes later.)



# What happens during Negotiation

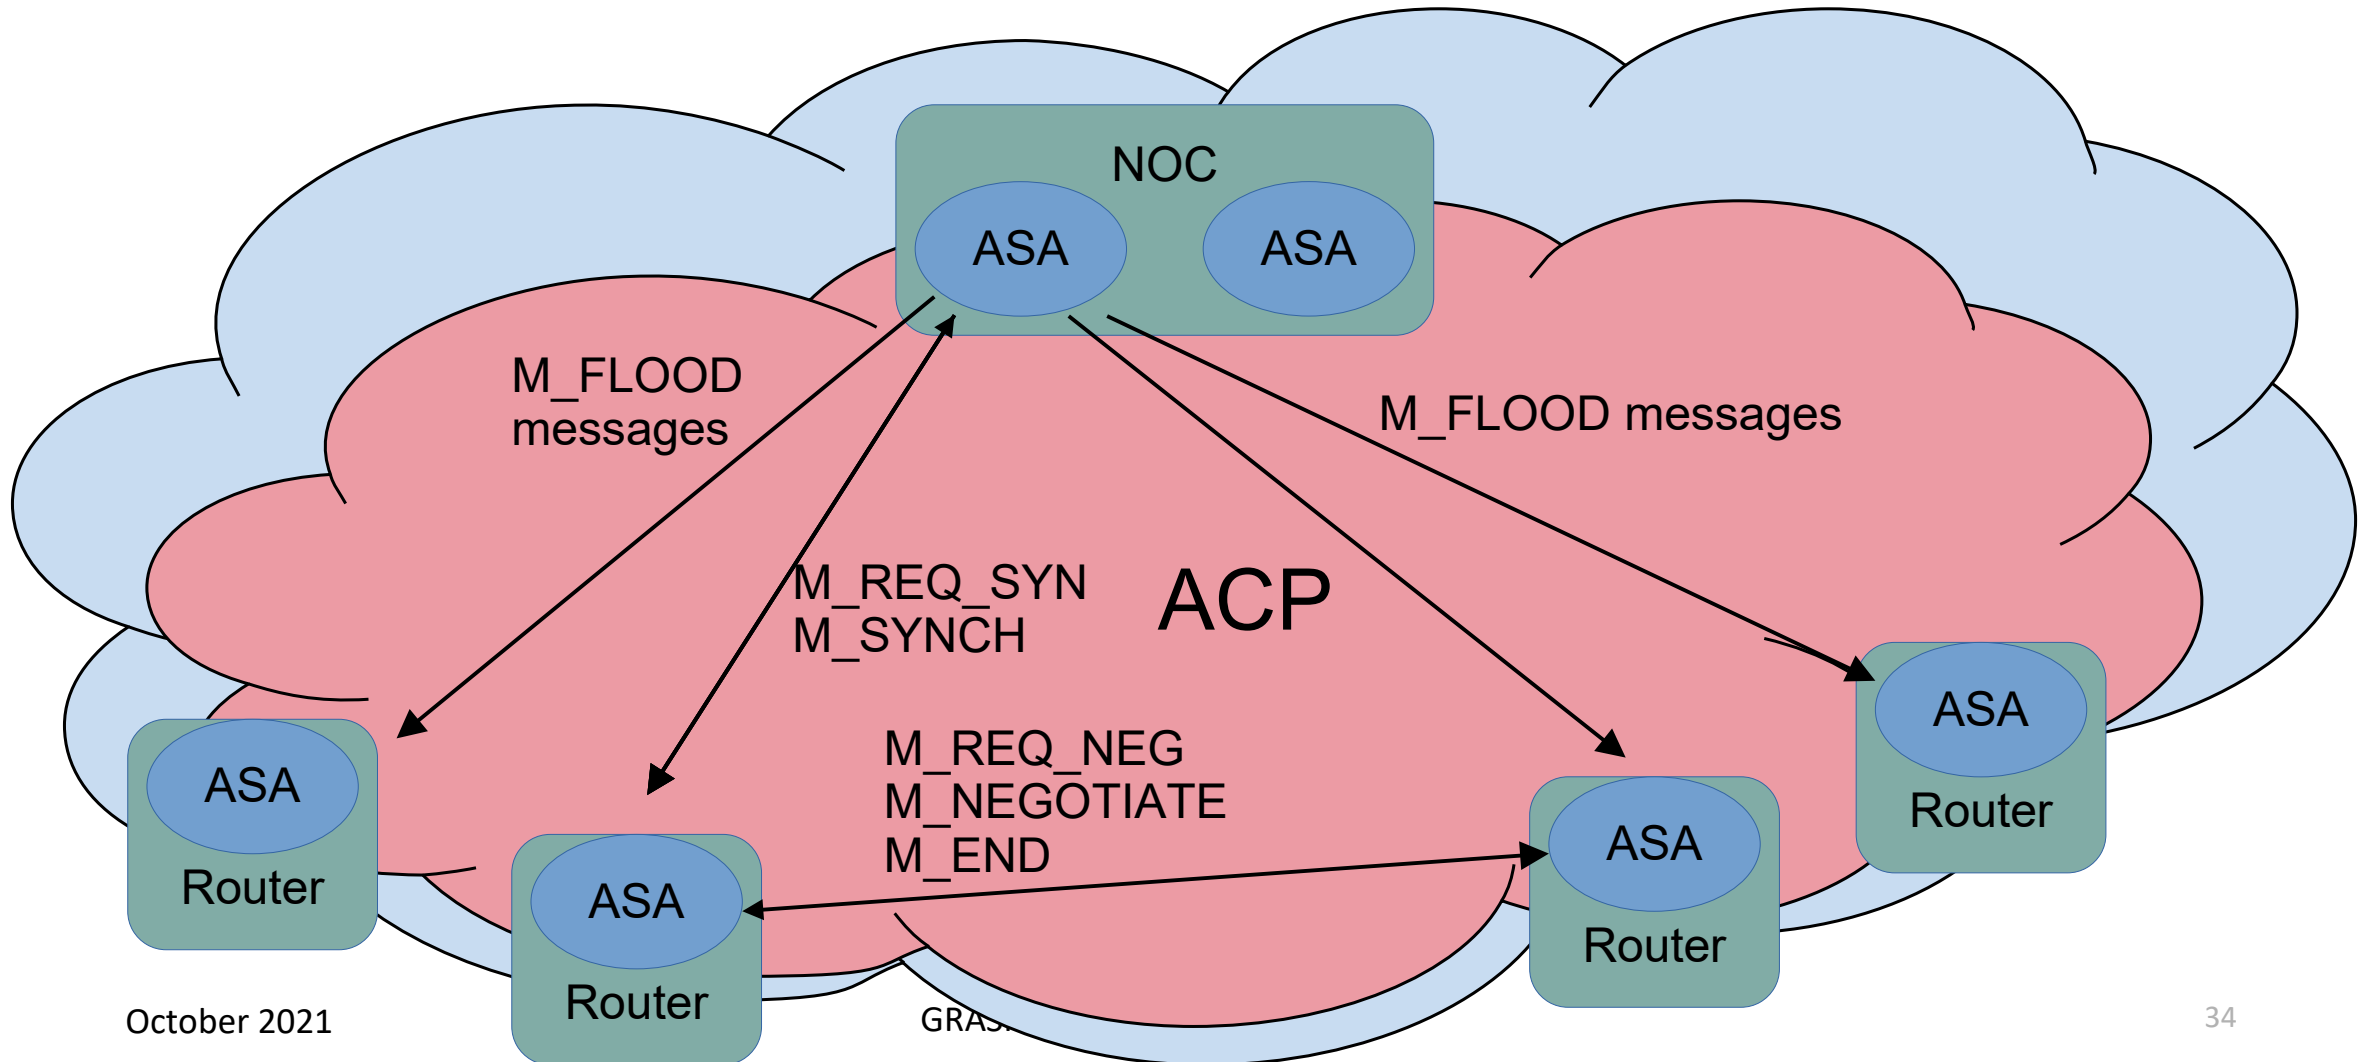
- An ASA able to negotiate a value for objective “A” waits for any incoming M\_REQ\_NEG message. This also makes the ASA discoverable for “A”.
  - This is server-like behavior
- An ASA needing to negotiate a value for “A” first uses discovery to find all peers that support “A” and chooses one of them. It then originates a unicast M\_NEGOTIATE to the chosen peer.
- The two peers swap alternate M\_NEGOTIATE messages containing a proffered value of “A” until one of them ends the negotiation with M\_END (carrying O\_ACCEPT or O\_DECLINE).
  - Note that a failed negotiation is not a protocol error.
- Either peer can insert an M\_WAIT message to delay the timeout.

(A diagram for negotiation comes later.)

# Interlude (7)



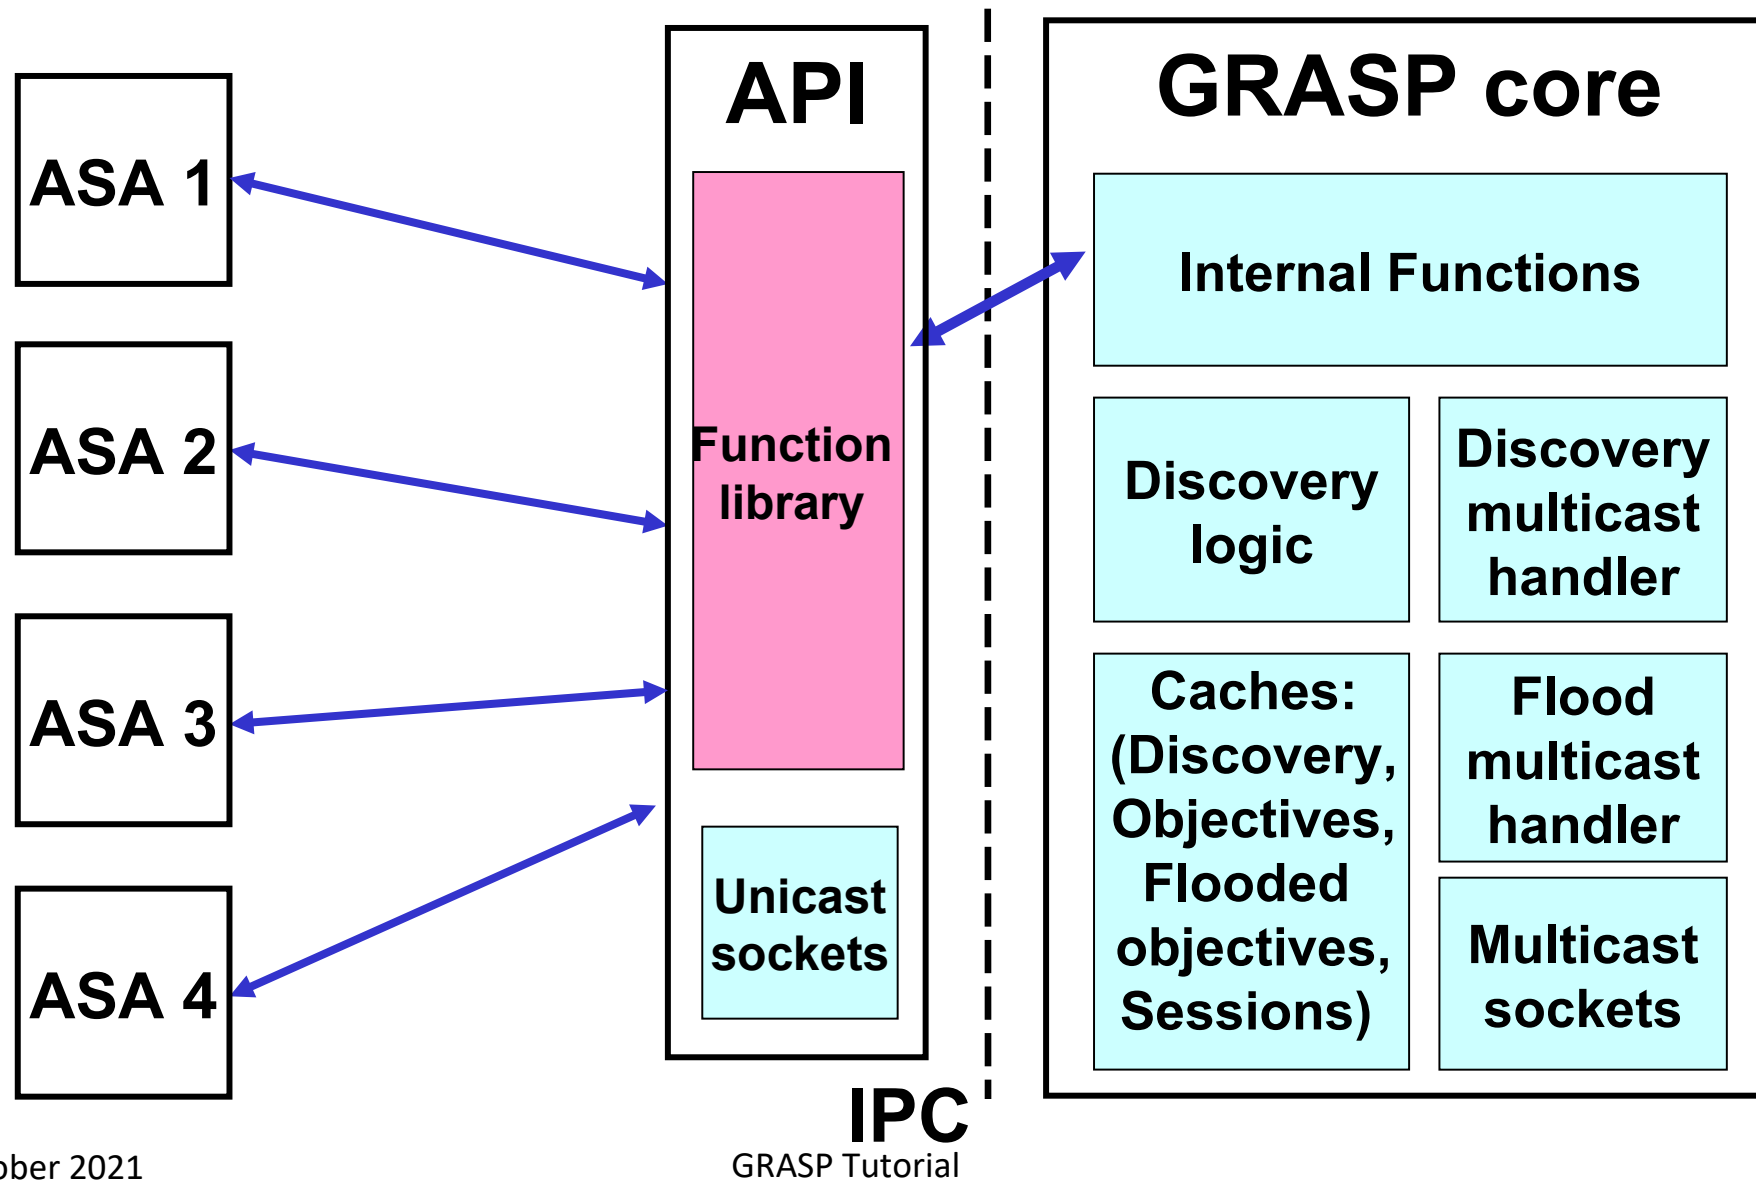
- Add objectives and messages



# Confused? Then you need an API.

- Although GRASP has few message types, it has quite powerful capabilities:
  - Discovery
  - Flooding
  - Client/server data synchronization
  - Peer-to-peer data negotiation
- Each of these has a clear purpose in autonomic functions.
- Each has its own complexity.
- In some cases, a subset of GRASP will be built into an application. Often, it will appear as an API. We will now discuss the ASA programmer's view of this API.

# ASA programmer's view



# GRASP API Functions

- *Registration*. An ASA can register itself and register the GRASP Objectives it manipulates.
- *Discovery*. An ASA can discover a peer willing to respond for a particular objective.
- *Negotiation*. An ASA can act as an initiator (requester) or responder (listener) for a negotiation session. Negotiation is a symmetric process, so most functions can be used by either party.
- *Synchronization*. An ASA can act as an initiator (requester) or responder (listener and data source) for data synchronization.
- *Flooding*. An ASA can send and receive a GRASP Objective that is flooded to all nodes of the ACP.

*The following slides show a simplified Python rendering of the main API functions.*

# Data types

**class objective(name)**

Attributes include **.loop\_count** and **.value**

**class asa\_locator()**

Attributes include **.locator** and **.is\_ipaddress**

**class tagged\_objective()**

Attributes **.objective** and **.locator**

# Registration Functions

## **register\_asa(asa\_name)**

Tells GRASP that a new ASA is starting.

Returns **(zero, asa\_handle)** if successful,  
**(errorcode, None)** if failure.

The ASA must use **asa\_handle** in every subsequent call.

## **register\_obj(asa\_handle, objective)**

Tells GRASP that the ASA will support the given objective.

Returns **(zero)** if successful,  
**(errorcode)** if failure.

# Discovery Function

**discover(asa\_handle, objective, timeout)**

Returns **(zero, list of asa\_locator)** if successful,  
**(errorcode, [])** if failure.



# Flooding Functions

## **flood(asa\_handle, ttl, tagged\_objectives)**

Floods a list of objectives with values to all nodes in the autonomic network.

Returns **zero** if successful,  
**errorcode** if failure.

## **get\_flood(asa\_handle, objective)**

Fetches flooded objectives from local cache.

Returns **(zero, tagged\_objectives)** if successful,  
**(errorcode, None)** if failure.

# Synchronization Functions

## **listen\_synchronize(*asa\_handle*, *objective*)**

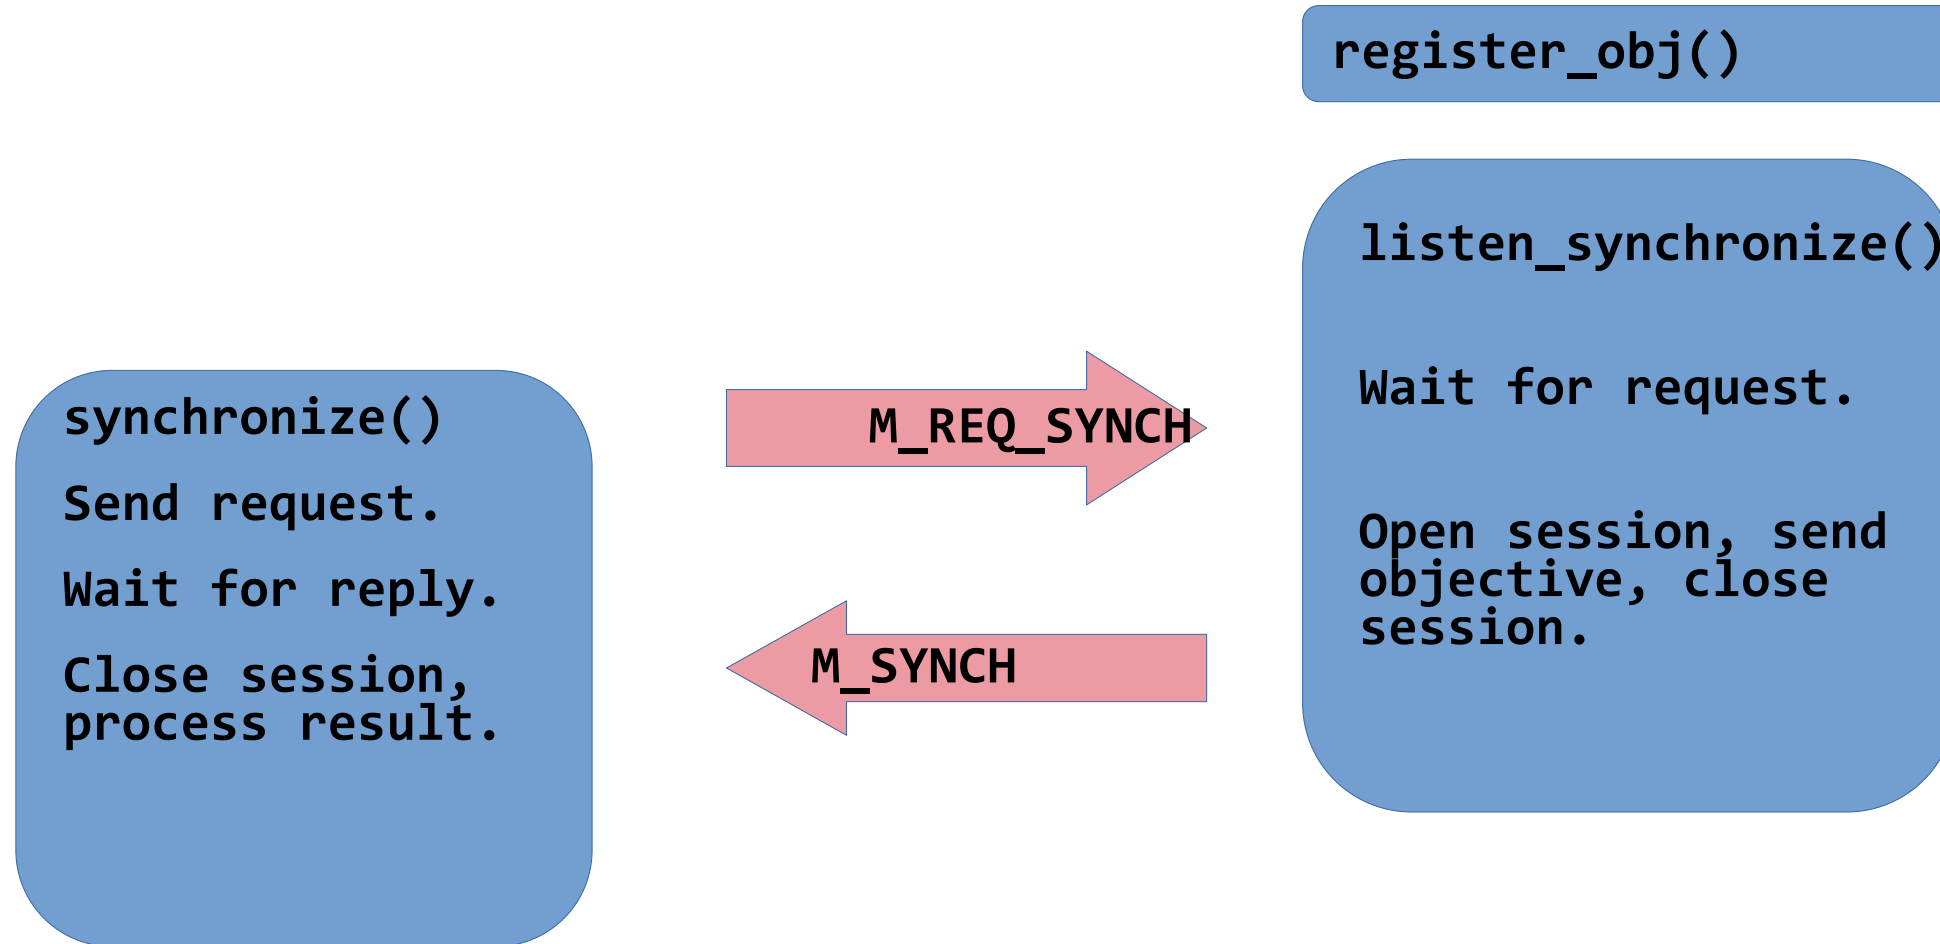
GRASP will listen for incoming synchronization requests and reply with the given objective and its value. Must be a separate thread.

## **synchronize(*asa\_handle*, *objective*, *locator*, *timeout*)**

Requests synchronized value of the given objective. The **locator** is an **asa\_locator** as returned by **discover()**.

Returns **(zero, objective)** if successful,  
**(errorcode, None)** if failure.

# A synchronization session



# Negotiation Functions (1)

**listen\_negotiate(asa\_handle, objective)**

Listen for incoming requests and start a negotiation session. Must be a separate thread.

Returns (**zero**, **session\_handle**, **requested\_objective**) if successful.

The **session\_handle** must be used in subsequent calls.

The value of the **requested\_objective** is the peer's initial offer for negotiation.

## Negotiation Functions (2)

**request\_negotiate(*asa\_handle*, *objective*, *peer*, *timeout*)**

Requests negotiation of the given objective, starting with its current value. The **peer** is an **asa\_locator** as returned by **discover()**.

There are 4 possible returns:

**(zero, None, objective, None)** The peer agreed with the offered value.

**(zero, session\_handle, objective, None)**

Negotiation continues. The returned objective contains the value offered by the peer.

**(errors.declined, None, None, string)**

The peer declined further negotiation, the string gives a reason

**(errorcode, None, None, None)** Some other error.

# Negotiation Functions (3)

After **request\_negotiate()** various functions are used symmetrically by either side:

**negotiate\_step(asa\_handle, session\_handle,  
                  objective, timeout)**

Sends the next proffered value of the objective to the peer. The returns are identical to **request\_negotiate()**. Used alternately by the two peers.

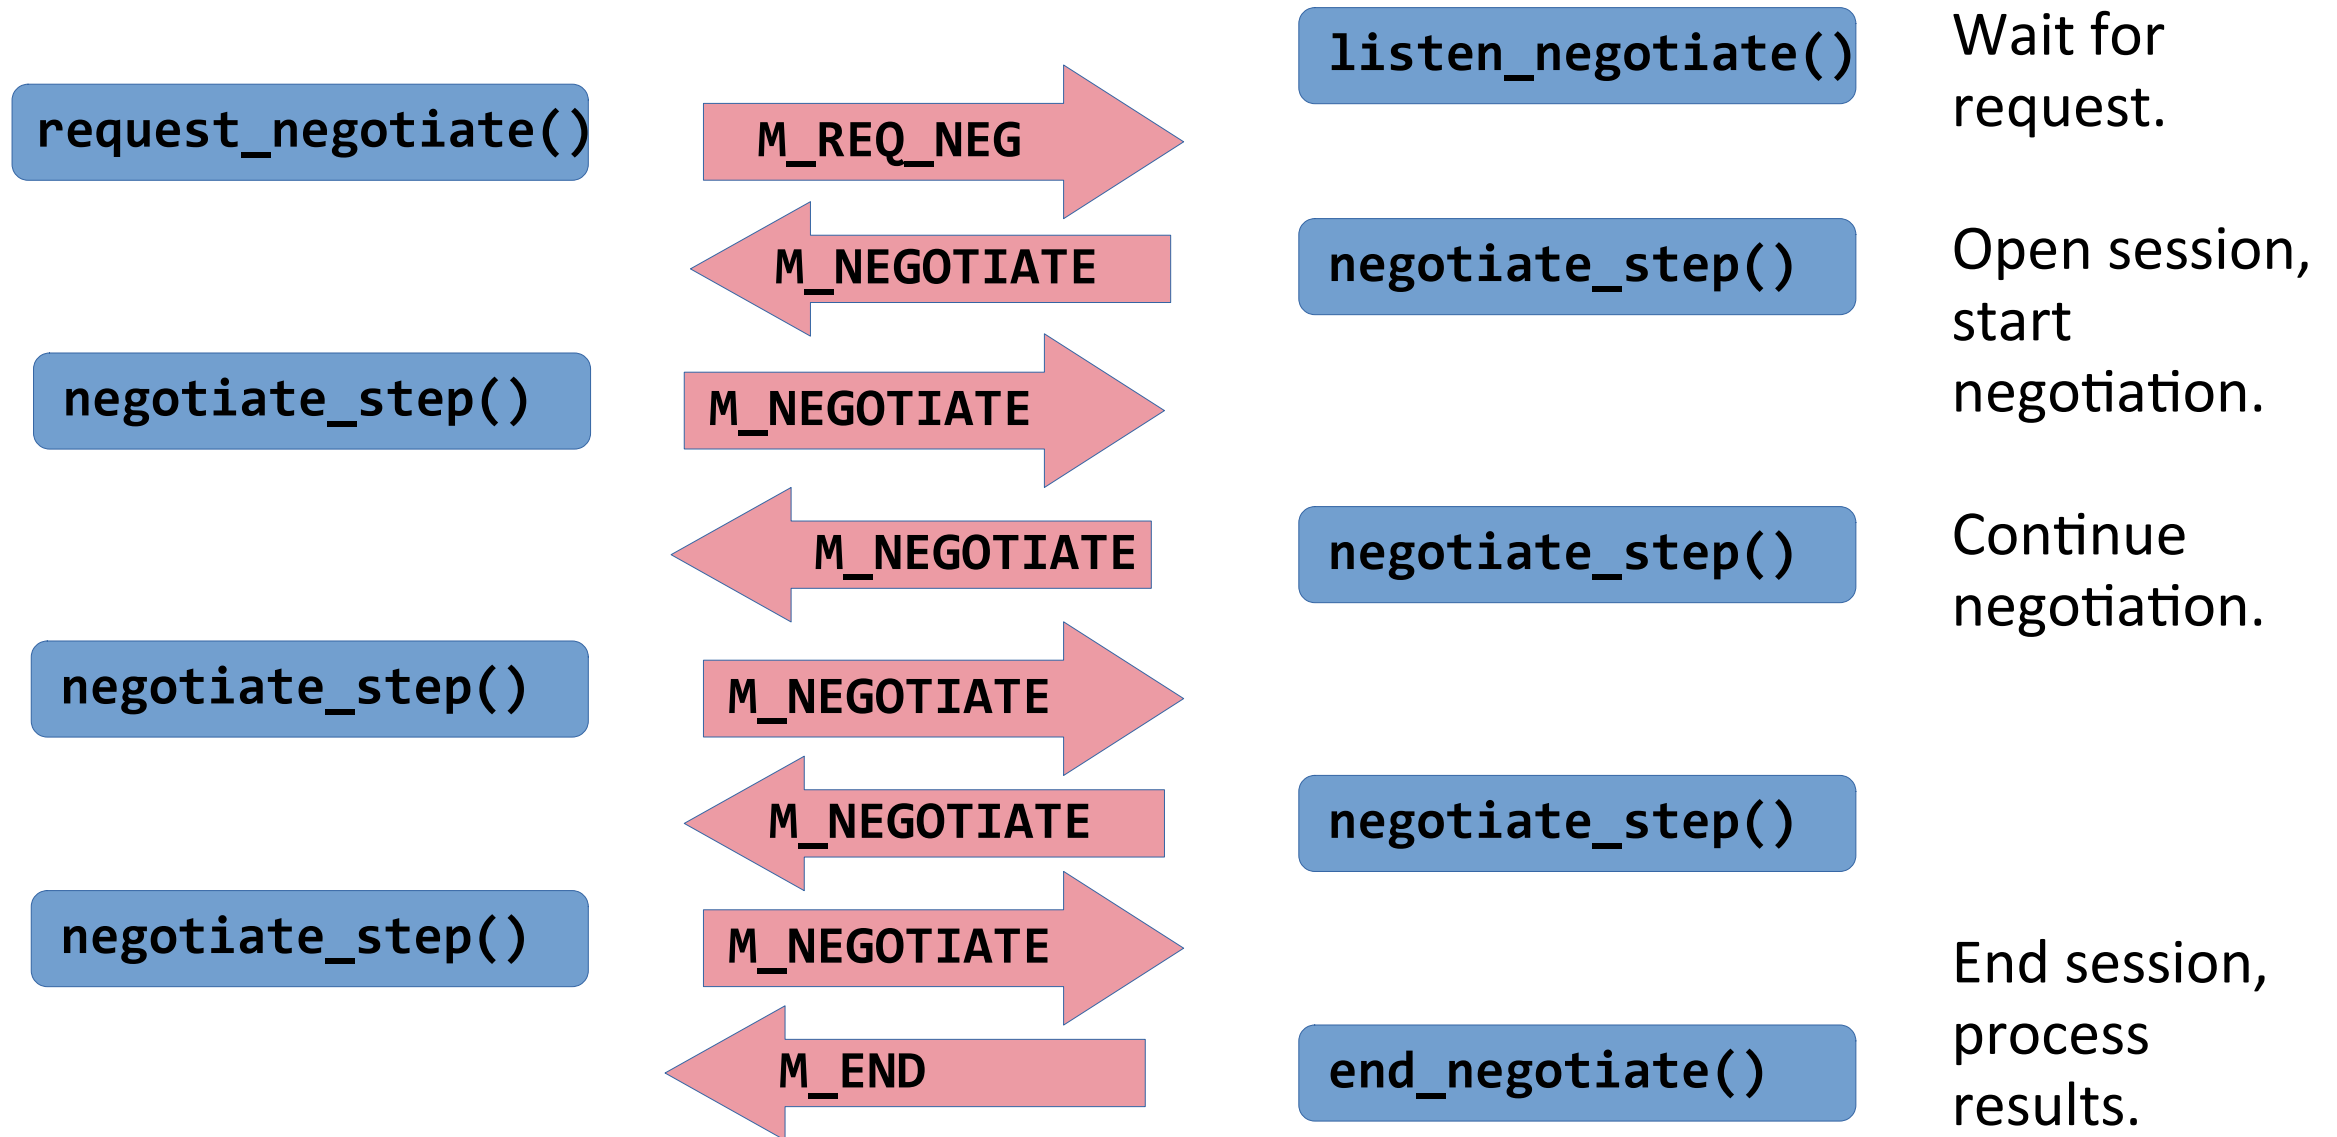
**negotiate\_wait(asa\_handle, session\_handle, timeout)**

Extend the timeout.

**end\_negotiate(asa\_handle, session\_handle, result)**

End negotiation (**result=True** for success, **False** to decline further negotiation.)

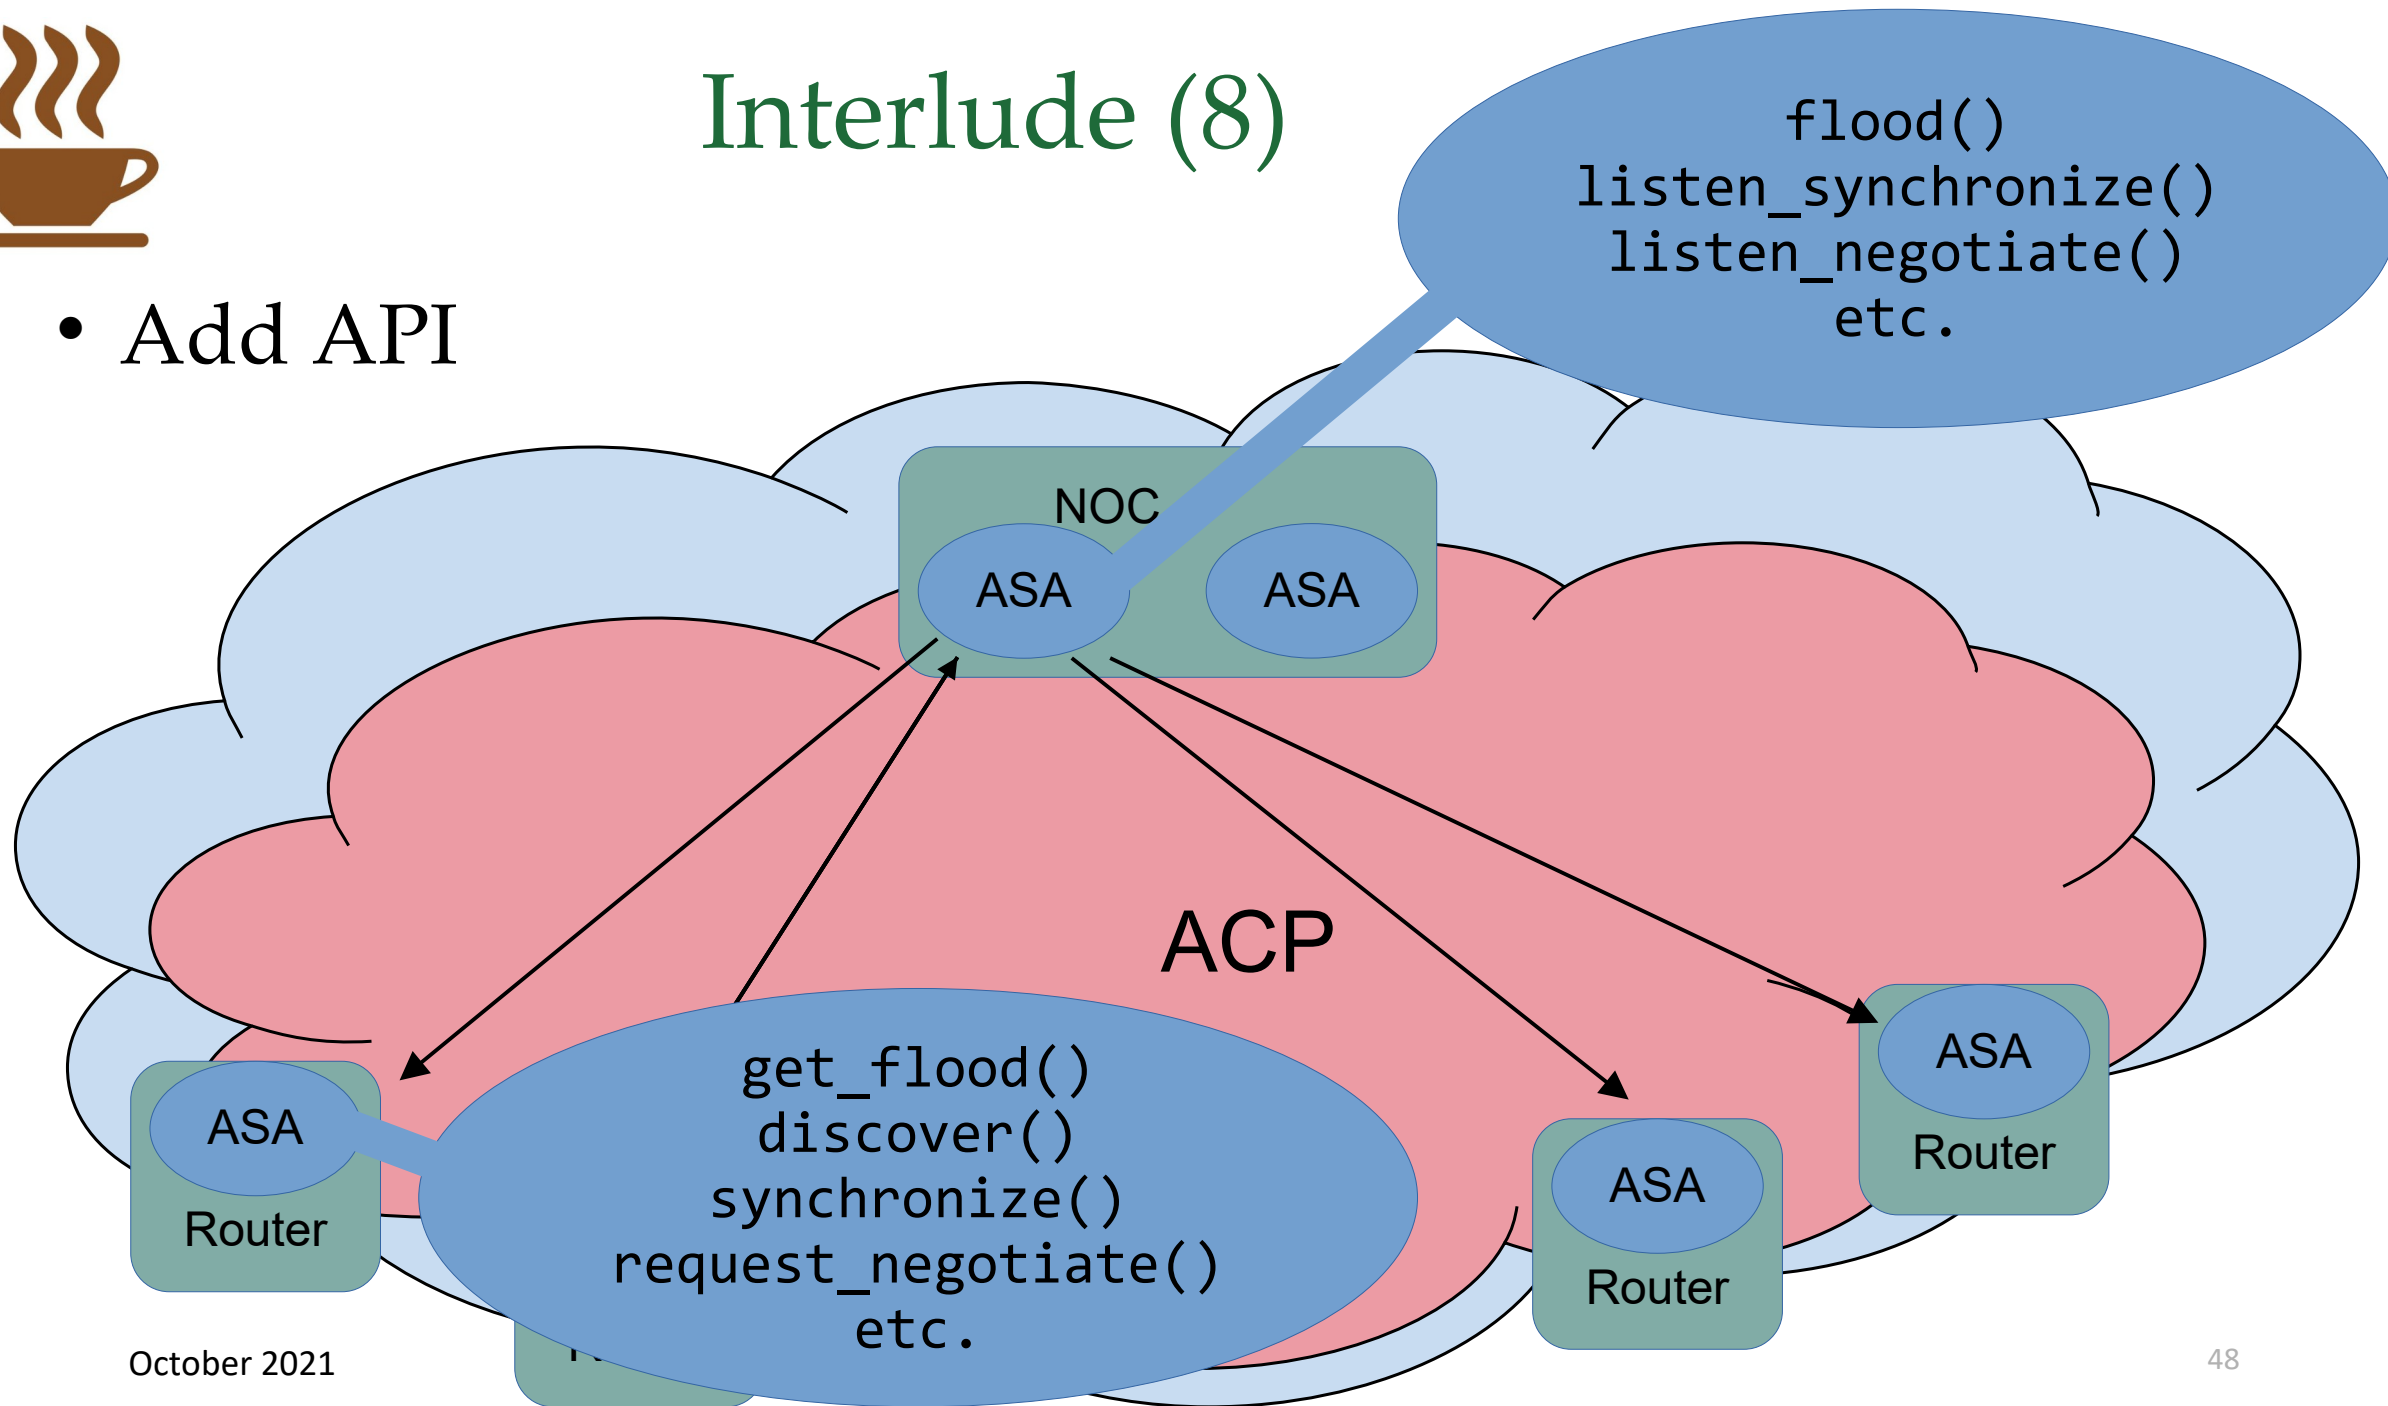
# A negotiation session





## Interlude (8)

- Add API





## Logic flows

- The API is designed for use in asynchronous operations such as handling multiple simultaneous negotiations, or processing floods, synchronizations, and negotiations at the same time.
- This could be done using an event-loop mechanism or a threading mechanism, depending on the programming environment in use.
- In the following slides, we show logic flows for an ASA that manages some (unnamed) distributed resource
  - We assume a threaded model
  - A very general outline is followed by pseudocode

# Logic flow outline – main thread

```
MAIN thread:  
initialise resource pool  
if origin:  
    start FLOODER to broadcast parameters  
start NEGOTIATOR and GARBAGE_COLLECTOR  
if not origin:  
    get resource parameters flooded by GRASP  
    start ASSIGN thread (allocates resources)  
do forever:  
    if resource pool is low:  
        negotiate for more resource from GRASP peer(s)
```

# Logic flow outline – other threads

**FLOODER** thread:

periodically flood resource parameters to all GRASP nodes

**NEGOTIATOR** thread:

wait for and satisfy negotiation requests from GRASP peers

**GARBAGE\_COLLECTOR** thread:

periodically compact the resource pool

**DELEGATOR** thread:

manage resource requests from non-autonomic devices & applications, assign resources from pool

# Pseudocode: MAIN (1)

# Initialization

Create empty resource\_pool

register\_asa()

register\_obj("EX1.Resource")

register\_obj("EX1.Params")

if origin:

    Obtain initial resource\_pool contents from NOC

    Obtain value of EX1.Params from NOC

    Start FLOODER thread to flood EX1.Params

    Start SYNCHRONIZER listener for EX1.Params

Start MAIN\_NEGOTIATOR thread for EX1.Resource

if not origin:

    get\_flood("EX1.Params")

    Start DELEGATOR thread

Start GARBAGE\_COLLECTOR thread

# Pseudocode: MAIN (2)

```
# main loop
do forever:
    if resource_pool is low:
        peers = discover("EX1.Resource")
        peer = #any choice among peers
        request_negotiate("EX1.Resource", peer)
        #Wait for response (M_NEGOTIATE, M_END or M_WAIT)
        if OK:
            if offered amount of resource sufficient:
                end_negotiate(True)
                Add resource to pool
                good_peer = peer
            else:
                end_negotiate(False) #negotiation failed
        sleep() #sleep time depends on application scenario
```

This is a very simple use of negotiation because it doesn't loop.

(A full negotiation needs an outer loop here.)

# Pseudocode: NEGOTIATOR

```
# MAIN_NEGOTIATOR thread:
do forever:
    listen_negotiate("EX1.Resource") # wait for M_REQ_NEG
    Start a separate new NEGOTIATOR thread for requested amount A

# NEGOTIATOR thread:
Request resource amount A from resource_pool
if not OK:
    while not OK and A > Amin:
        A = A-1
        Request resource amount A from resource_pool
if OK:
    negotiate_step("EX1.Resource") # Offer resource amount A
    if received M_END + O_ACCEPT:
        # negotiation succeeded
    elif received M_END + O_DECLINE or other error:
        # negotiation failed
else:
    end_negotiate(False) # negotiation failed
```

Will offer the best  
it can get from  
the resource  
pool.

Again, a single  
step  
negotiation

# Pseudocode: DELEGATOR

```
# There are no GRASP calls. This is actual resource assignment.
do forever:
    Wait for request or release for resource amount A
    if request:
        Get resource amount A from resource_pool
        if OK:
            Delegate resource to consumer
            Record in delegated_list
        else:
            Signal failure to consumer
            Signal main thread that resource_pool is low
    else:
        Delete resource from delegated_list
        Return resource amount A to resource_pool
```

# Pseudocode: other threads

**# SYNCHRONIZER thread:**

**do forever:**

**listen\_synchronize("EX1.Params")**

**# FLOODER thread:**

**do forever:**

**flood("EX1.Params")**

**sleep() #sleep time depends on application scenario**

**# GARBAGE\_COLLECTOR thread:**

**do forever:**

**Search resource\_pool for adjacent resources**

**Merge adjacent resources**

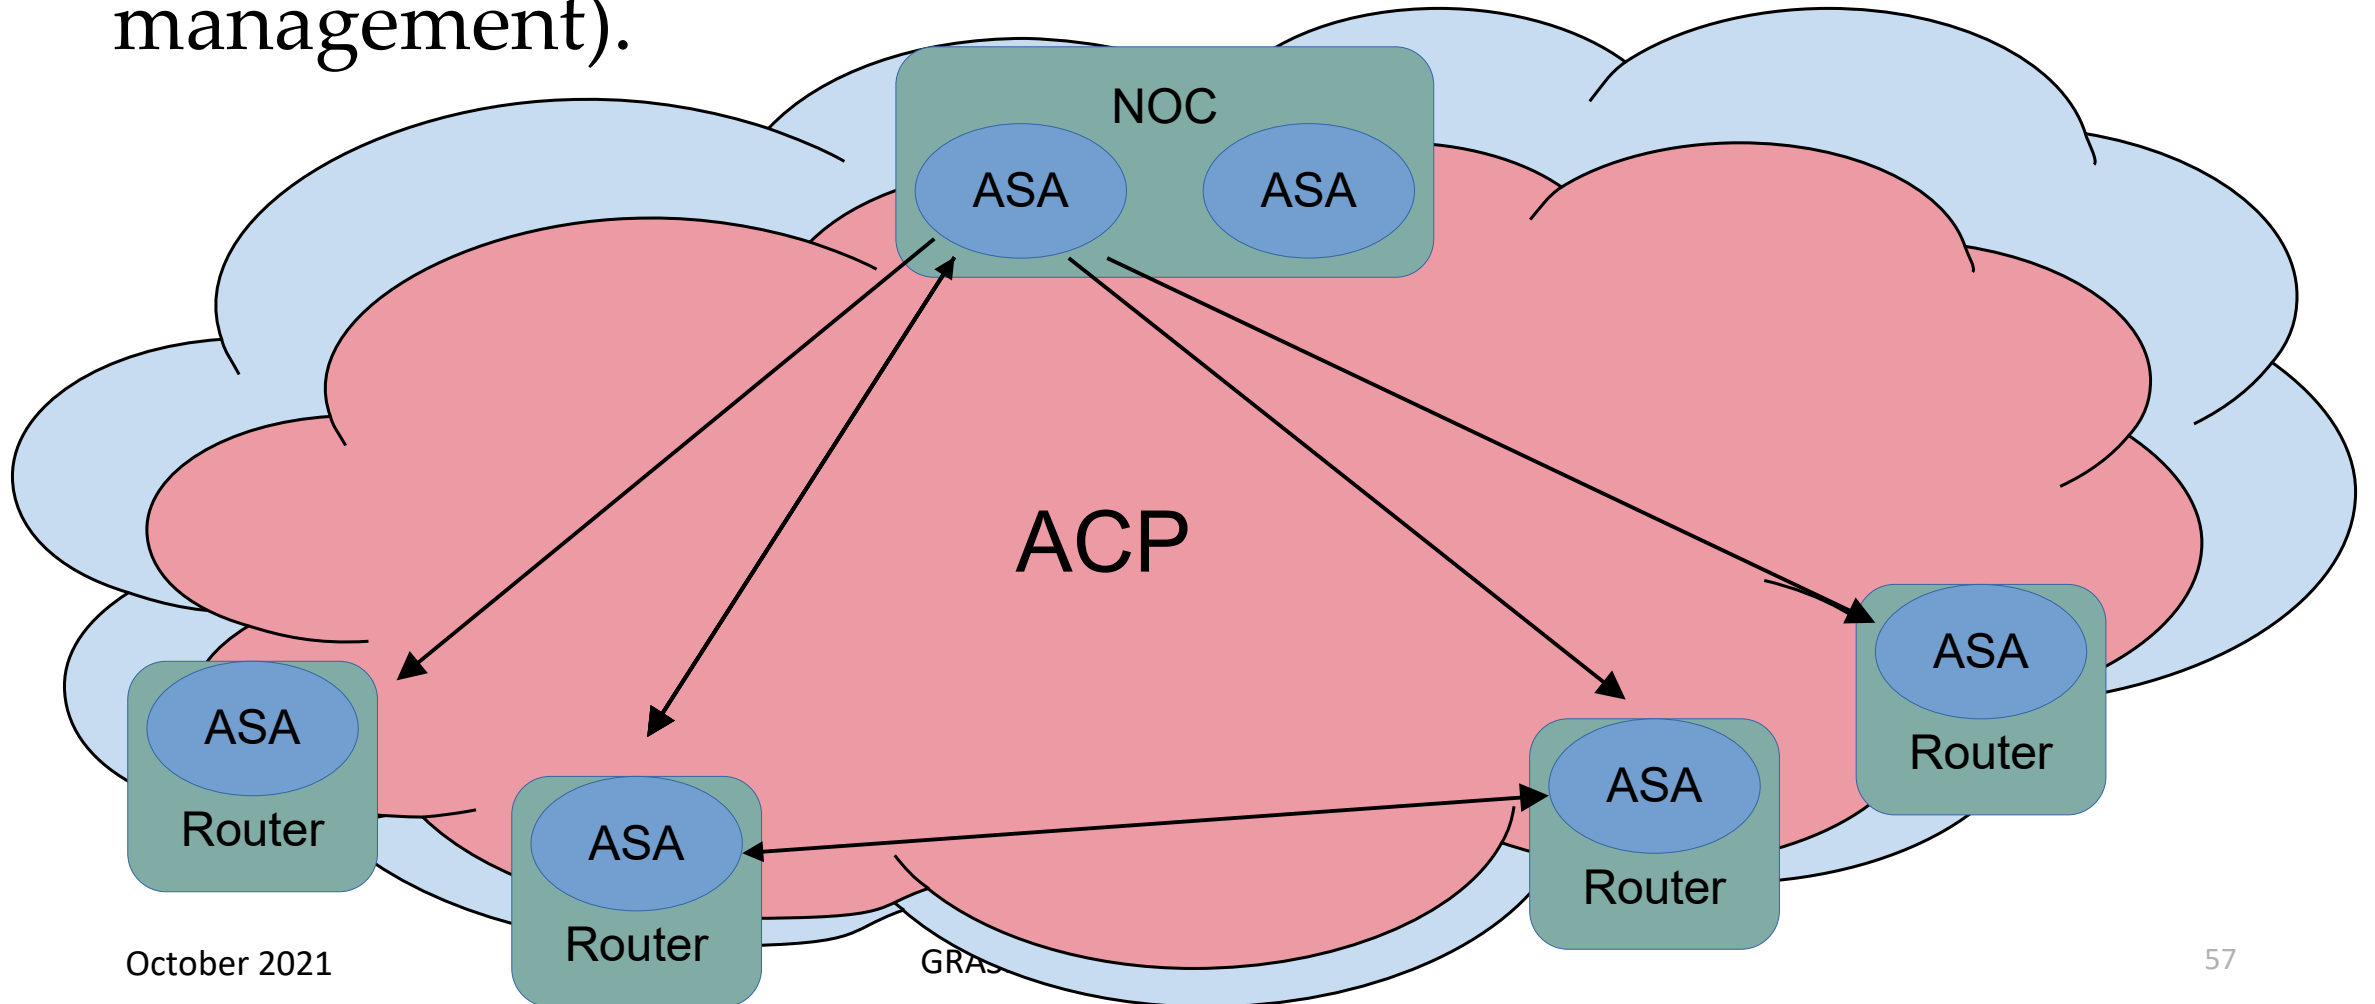
**sleep() #sleep time depends on application scenario**



# Final Interlude



- With those logic flows, we have designed the ASAs for a distributed resource manager (such as IP prefix management).



# Security

- GRASP does not have its own security mechanism. It is used over a secure and encrypted Autonomic Control Plane.
- TLS is recommended for the unicast messages.
- GRASP includes a very restricted subset for use during bootstrapping of the ACP, known as “Discovery Unsolicited Link-Local” (DULL).

## GRASP Prototype

- A Python 3 implementation of GRASP and its API as **graspi.py**
- About 2600 lines of code
- A test suite to exercise as many code paths as possible
- Various toy ASAs to test "real" operation across the network
  - bank/client negotiation
  - model of secure bootstrap process
  - model of IPv6 prefix management
  - bulk transfer using GRASP

<https://github.com/becarpenter/graspy>

Start with

<https://github.com/becarpenter/graspy/blob/master/graspy.pdf>

# The End

- Raise issues, questions and comments at <https://github.com/becarpenter/grasdoc/issues>
- Join the ANIMA WG at <https://www.ietf.org/mailman/listinfo/anima>
- Want to improve the Python prototype? It's open source on GitHub at <https://github.com/becarpenter/graspy>
- Got your own GRASP implementation? List it at <https://brski.org/grasp-impls.html>